

1 k_nearest_neighbor.py

```
1 from builtins import range
2 from builtins import object
3 import numpy as np
4 from past.builtins import xrange
5
6
7 class KNearestNeighbor(object):
8     """ a kNN classifier with L2 distance """
9
10    def __init__(self):
11        pass
12
13    def train(self, X, y):
14        """
15        Train the classifier. For k-nearest neighbors this is just
16        memorizing the training data.
17
18        Inputs:
19        - X: A numpy array of shape (num_train, D) containing the training data
20            consisting of num_train samples each of dimension D.
21        - y: A numpy array of shape (N,) containing the training labels, where
22            y[i] is the label for X[i].
23        """
24        self.X_train = X
25        self.y_train = y
26
27    def predict(self, X, k=1, num_loops=0):
28        """
29        Predict labels for test data using this classifier.
30
31        Inputs:
32        - X: A numpy array of shape (num_test, D) containing test data consisting
33            of num_test samples each of dimension D.
34        - k: The number of nearest neighbors that vote for the predicted labels.
35        - num_loops: Determines which implementation to use to compute distances
36            between training points and testing points.
37
38        Returns:
39        - y: A numpy array of shape (num_test,) containing predicted labels for the
40            test data, where y[i] is the predicted label for the test point X[i].
41        """
42        if num_loops == 0:
43            dists = self.compute_distances_no_loops(X)
44        elif num_loops == 1:
45            dists = self.compute_distances_one_loop(X)
46        elif num_loops == 2:
47            dists = self.compute_distances_two_loops(X)
48        else:
49            raise ValueError('Invalid value %d for num_loops' % num_loops)
50
51        return self.predict_labels(dists, k=k)
52
53    def compute_distances_two_loops(self, X):
54        """
55        Compute the distance between each test point in X and each training point
56        in self.X_train using a nested loop over both the training data and the
57        test data.
58
59        Inputs:
60        - X: A numpy array of shape (num_test, D) containing test data.
61
62        Returns:
63        - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
64            is the Euclidean distance between the ith test point and the jth training
65            point.
66        """
67        num_test = X.shape[0]
68        num_train = self.X_train.shape[0]
69        dists = np.zeros((num_test, num_train))
70        for i in range(num_test):
71            for j in range(num_train):
72                #####
73                # TODO:
74                #####
```

```

74         # Compute the l2 distance between the ith test point and the jth #
75         # training point, and store the result in dists[i, j]. You should #
76         # not use a loop over dimension, nor use np.linalg.norm(). #
77         #####
78         # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
79
80         dists[i, j] = np.sqrt(np.sum(np.square(X[i] - self.X_train[j])))
81
82         # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
83     return dists
84
85 def compute_distances_one_loop(self, X):
86     """
87     Compute the distance between each test point in X and each training point
88     in self.X_train using a single loop over the test data.
89
90     Input / Output: Same as compute_distances_two_loops
91     """
92     num_test = X.shape[0]
93     num_train = self.X_train.shape[0]
94     dists = np.zeros((num_test, num_train))
95     for i in range(num_test):
96         #####
97         # TODO: #
98         # Compute the l2 distance between the ith test point and all training #
99         # points, and store the result in dists[i, :]. #
100        # Do not use np.linalg.norm(). #
101        #####
102        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
103
104        dists[i] = np.sqrt(np.sum(np.square(X[i] - self.X_train), axis = 1))
105
106        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
107    return dists
108
109 def compute_distances_no_loops(self, X):
110     """
111     Compute the distance between each test point in X and each training point
112     in self.X_train using no explicit loops.
113
114     Input / Output: Same as compute_distances_two_loops
115     """
116     num_test = X.shape[0]
117     num_train = self.X_train.shape[0]
118     dists = np.zeros((num_test, num_train))
119     #####
120     # TODO: #
121     # Compute the l2 distance between all test points and all training #
122     # points without using any explicit loops, and store the result in #
123     # dists. #
124     # #
125     # You should implement this function using only basic array operations; #
126     # in particular you should not use functions from scipy, #
127     # nor use np.linalg.norm(). #
128     # #
129     # HINT: Try to formulate the l2 distance using matrix multiplication #
130     # and two broadcast sums. #
131     #####
132     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
133
134     testSumOfSquare = np.sum(np.square(X), axis=1)
135     # Get a (500,) matrix
136     trainSumOfSquare = np.sum(np.square(self.X_train), axis=1)
137     # Get a (5000,) matrix
138     product = np.dot(X, self.X_train.T)
139     # Get a (500, 5000) matrix
140     dists = np.sqrt(testSumOfSquare[:, np.newaxis] + trainSumOfSquare - 2*product)
141     # Add a new axis to testSumOfSquare to create a (500, 1) matrix and plus a (5000,) matrix to get a
142     # (500, 5000) matrix
143
144     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
145     return dists
146
147 def predict_labels(self, dists, k=1):
148     """
149     Given a matrix of distances between test points and training points,

```

```

149     predict a label for each test point.
150
151     Inputs:
152     - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
153       gives the distance between the ith test point and the jth training point.
154
155     Returns:
156     - y: A numpy array of shape (num_test,) containing predicted labels for the
157       test data, where y[i] is the predicted label for the test point X[i].
158     """
159     num_test = dists.shape[0]
160     y_pred = np.zeros(num_test)
161     for i in range(num_test):
162         # A list of length k storing the labels of the k nearest neighbors to
163         # the ith test point.
164         closest_y = []
165         #####
166         # TODO:
167         # Use the distance matrix to find the k nearest neighbors of the ith
168         # testing point, and use self.y_train to find the labels of these
169         # neighbors. Store these labels in closest_y.
170         # Hint: Look up the function numpy.argsort.
171         #####
172         # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
173
174         index_ascending = np.argsort(dists[i])
175         closest_y = self.y_train[index_ascending[:k]]
176
177
178         # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
179         #####
180         # TODO:
181         # Now that you have found the labels of the k nearest neighbors, you
182         # need to find the most common label in the list closest_y of labels.
183         # Store this label in y_pred[i]. Break ties by choosing the smaller
184         # label.
185         #####
186         # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
187
188         y_pred[i] = np.argmax(np.bincount(closest_y))
189
190         # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
191
192     return y_pred

```

2 linear_classifier.py

```
1 from __future__ import print_function
2
3 from builtins import range
4 from builtins import object
5 import numpy as np
6 from cs231n.classifiers.linear_svm import *
7 from cs231n.classifiers.softmax import *
8 from past.builtins import xrange
9
10
11 class LinearClassifier(object):
12
13     def __init__(self):
14         self.W = None
15
16     def train(self, X, y, learning_rate=1e-3, reg=1e-5, num_iters=100,
17             batch_size=200, verbose=False):
18         """
19         Train this linear classifier using stochastic gradient descent.
20
21         Inputs:
22         - X: A numpy array of shape (N, D) containing training data; there are N
23             training samples each of dimension D.
24         - y: A numpy array of shape (N,) containing training labels; y[i] = c
25             means that X[i] has label 0 ≤ c < C for C classes.
26         - learning_rate: (float) learning rate for optimization.
27         - reg: (float) regularization strength.
28         - num_iters: (integer) number of steps to take when optimizing
29         - batch_size: (integer) number of training examples to use at each step.
30         - verbose: (boolean) If true, print progress during optimization.
31
32         Outputs:
33         A list containing the value of the loss function at each training iteration.
34         """
35         num_train, dim = X.shape
36         num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
37         if self.W is None:
38             # lazily initialize W
39             self.W = 0.001 * np.random.randn(dim, num_classes)
40
41         # Run stochastic gradient descent to optimize W
42         loss_history = []
43         for it in range(num_iters):
44             X_batch = None
45             y_batch = None
46
47             #####
48             # TODO:
49             # Sample batch_size elements from the training data and their
50             # corresponding labels to use in this round of gradient descent.
51             # Store the data in X_batch and their corresponding labels in
52             # y_batch; after sampling X_batch should have shape (batch_size, dim)
53             # and y_batch should have shape (batch_size,)
54             #
55             # Hint: Use np.random.choice to generate indices. Sampling with
56             # replacement is faster than sampling without replacement.
57             #####
58             # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
59
60             indices_of_batch = np.random.choice(num_train, batch_size)
61             X_batch = X[indices_of_batch]
62             y_batch = y[indices_of_batch]
63
64             # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
65
66             # evaluate loss and gradient
67             loss, grad = self.loss(X_batch, y_batch, reg)
68             loss_history.append(loss)
69
70             # perform parameter update
71             #####
72             # TODO:
73             # Update the weights using the gradient and the learning rate.
74             #
```

```

74 #####
75 # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
76
77     self.W -= learning_rate * grad
78
79 # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
80
81     if verbose and it % 100 == 0:
82         print('iteration %d / %d: loss %f' % (it, num_iters, loss))
83
84     return loss_history
85
86 def predict(self, X):
87     """
88     Use the trained weights of this linear classifier to predict labels for
89     data points.
90
91     Inputs:
92     - X: A numpy array of shape (N, D) containing training data; there are N
93         training samples each of dimension D.
94
95     Returns:
96     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
97         array of length N, and each element is an integer giving the predicted
98         class.
99     """
100     y_pred = np.zeros(X.shape[0])
101     #####
102     # TODO:
103     # Implement this method. Store the predicted labels in y_pred.
104     #####
105     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
106
107     # Get the largest index of every row in X.dot(self.W)
108     y_pred = np.argmax(X.dot(self.W), axis = 1)
109
110     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
111     return y_pred
112
113 def loss(self, X_batch, y_batch, reg):
114     """
115     Compute the loss function and its derivative.
116     Subclasses will override this.
117
118     Inputs:
119     - X_batch: A numpy array of shape (N, D) containing a minibatch of N
120         data points; each point has dimension D.
121     - y_batch: A numpy array of shape (N,) containing labels for the minibatch.
122     - reg: (float) regularization strength.
123
124     Returns: A tuple containing:
125     - loss as a single float
126     - gradient with respect to self.W; an array of the same shape as W
127     """
128     pass
129
130
131 class LinearSVM(LinearClassifier):
132     """ A subclass that uses the Multiclass SVM loss function """
133
134     def loss(self, X_batch, y_batch, reg):
135         return svm_loss_vectorized(self.W, X_batch, y_batch, reg)
136
137
138 class Softmax(LinearClassifier):
139     """ A subclass that uses the Softmax + Cross-entropy loss function """
140
141     def loss(self, X_batch, y_batch, reg):
142         return softmax_loss_vectorized(self.W, X_batch, y_batch, reg)

```

3 linear_svm.py

```
1 from builtins import range
2 import numpy as np
3 from random import shuffle
4 from past.builtins import xrange
5
6 def svm_loss_naive(W, X, y, reg):
7     """
8     Structured SVM loss function, naive implementation (with loops).
9
10    Inputs have dimension D, there are C classes, and we operate on minibatches
11    of N examples.
12
13    Inputs:
14    - W: A numpy array of shape (D, C) containing weights.
15    - X: A numpy array of shape (N, D) containing a minibatch of data.
16    - y: A numpy array of shape (N,) containing training labels; y[i] = c means
17        that X[i] has label c, where 0 <= c < C.
18    - reg: (float) regularization strength
19
20    Returns a tuple of:
21    - loss as single float
22    - gradient with respect to weights W; an array of same shape as W
23    """
24    dW = np.zeros(W.shape) # initialize the gradient as zero
25
26    # compute the loss and the gradient
27    num_classes = W.shape[1]
28    num_train = X.shape[0]
29    loss = 0.0
30    for i in range(num_train):
31        scores = X[i].dot(W)
32        correct_class_score = scores[y[i]]
33        for j in range(num_classes):
34            if j == y[i]:
35                continue
36            margin = scores[j] - correct_class_score + 1 # note delta = 1
37            if margin > 0:
38                loss += margin
39
40        # Subtract X[i].T from the y[i]-th column of dW
41        dW[:, y[i]] -= X[i].T
42
43        # Add X[i].T to the j-th column of dW if j != y[i]
44        dW[:, j] += X[i].T
45
46    # Right now the loss is a sum over all training examples, but we want it
47    # to be an average instead so we divide by num_train.
48    loss /= num_train
49
50    # Add regularization to the loss.
51    loss += reg * np.sum(W * W)
52
53    #####
54    # TODO:
55    # Compute the gradient of the loss function and store it dW.
56    # Rather than first computing the loss and then computing the derivative,
57    # it may be simpler to compute the derivative at the same time that the
58    # loss is being computed. As a result you may need to modify some of the
59    # code above to compute the gradient.
60    #####
61    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
62
63    # Want dW to be an average, divide it by num_train
64    dW /= num_train
65
66    # Add regularization to the gradient, which is the derivative of loss with respect to W
67    dW += 2 * reg * W
68
69    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
70
71    return loss, dW
72
73
```

```

74
75 def svm_loss_vectorized(W, X, y, reg):
76     """
77     Structured SVM loss function, vectorized implementation.
78
79     Inputs and outputs are the same as svm_loss_naive.
80     """
81     loss = 0.0
82     dW = np.zeros(W.shape) # initialize the gradient as zero
83
84     #####
85     # TODO:
86     # Implement a vectorized version of the structured SVM loss, storing the
87     # result in loss.
88     #####
89     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
90     num_classes = W.shape[1]
91     num_train = X.shape[0]
92     scores = X.dot(W)
93
94     # Diagonal elements are correct scores. Get the correct scores for every class, it is a (N,) matrix.
95     correct_scores = scores[np.arange(num_train), y]
96
97     # Clone the correct_scores column num_classes times to get a (N, C) matrix.
98     correct_scores_matrix = np.array([correct_scores,] * num_classes).transpose()
99
100    # Use scores to subtract correct scores.
101    margins = np.maximum(0, scores - correct_scores_matrix + 1)
102
103    # Diagonal elements in margins should be 0.
104    margins[np.arange(num_train), y] = 0
105
106    loss = np.sum(margins)
107    loss /= num_train
108    loss += reg * np.sum(W * W)
109
110    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
111
112    #####
113    # TODO:
114    # Implement a vectorized version of the gradient for the structured SVM
115    # loss, storing the result in dW.
116    #
117    # Hint: Instead of computing the gradient from scratch, it may be easier
118    # to reuse some of the intermediate values that you used to compute the
119    # loss.
120    #####
121    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
122
123    # Create a margins_for_dW which has the same shape as margins
124    margins_for_dW = np.zeros(margins.shape)
125
126    # If there is a positive value in margins, the corresponding element in margins_for_dW is 1.
127    margins_for_dW[margins > 0] = 1
128
129    # Count the number of positive values in margins
130    sum_of_row = np.sum(margins_for_dW, axis = 1)
131
132    # The diagonal elements in margins_for_dW should be the opposite value of sum_of_row.
133    margins_for_dW[np.arange(num_train), y] = -sum_of_row
134
135    # Use X.T multiply sum_of_row
136    dW = X.T.dot(margins_for_dW)/num_train + 2*reg*W
137
138    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
139
140    return loss, dW

```

4 softmax.py

```
1 from builtins import range
2 import numpy as np
3 from random import shuffle
4 from past.builtins import xrange
5
6 def softmax_loss_naive(W, X, y, reg):
7     """
8     Softmax loss function, naive implementation (with loops)
9
10    Inputs have dimension D, there are C classes, and we operate on minibatches
11    of N examples.
12
13    Inputs:
14    - W: A numpy array of shape (D, C) containing weights.
15    - X: A numpy array of shape (N, D) containing a minibatch of data.
16    - y: A numpy array of shape (N,) containing training labels; y[i] = c means
17        that X[i] has label c, where 0 <= c < C.
18    - reg: (float) regularization strength
19
20    Returns a tuple of:
21    - loss as single float
22    - gradient with respect to weights W; an array of same shape as W
23    """
24    # Initialize the loss and gradient to zero.
25    loss = 0.0
26    dW = np.zeros_like(W)
27
28    #####
29    # TODO: Compute the softmax loss and its gradient using explicit loops. #
30    # Store the loss in loss and the gradient in dW. If you are not careful #
31    # here, it is easy to run into numeric instability. Don't forget the #
32    # regularization! #
33    #####
34    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
35
36    num_classes = W.shape[1]
37    num_train = X.shape[0]
38    for i in range(num_train):
39        scores = X[i].dot(W)
40
41        # In order to avoid numeric instability, make the highest score to be zero
42        scores -= np.max(scores)
43
44        scores_exp_sum = np.sum(np.exp(scores))
45        correct_score_exp = np.exp(scores[y[i]])
46        loss += -np.log(correct_score_exp / scores_exp_sum)
47
48        # Calculate the gradient of W
49        for j in range(num_classes):
50            dW[:, j] += - ((j == y[i]) - (np.exp(scores[j]) / scores_exp_sum)) * X[i]
51
52    loss /= num_train
53    loss += reg * np.sum(W * W)
54    dW /= num_train
55    dW += 2 * reg * W
56
57
58    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
59
60    return loss, dW
61
62
63
64 def softmax_loss_vectorized(W, X, y, reg):
65     """
66     Softmax loss function, vectorized version.
67
68     Inputs and outputs are the same as softmax_loss_naive.
69     """
70    # Initialize the loss and gradient to zero.
71    loss = 0.0
72    dW = np.zeros_like(W)
73
```



```

74 #####
75 # TODO: Compute the softmax loss and its gradient using no explicit loops. #
76 # Store the loss in loss and the gradient in dW. If you are not careful #
77 # here, it is easy to run into numeric instability. Don't forget the #
78 # regularization! #
79 #####
80 # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
81
82 num_classes = W.shape[1]
83 num_train = X.shape[0]
84 scores = X.dot(W)
85
86 # Get the maximum score of every row and clone it multiple times to be a (N, C) matrix
87 max_scores = np.amax(scores, axis = 1)
88 max_scores_matrix = np.array([max_scores,] * num_classes).transpose()
89 scores -= max_scores_matrix
90
91 correct_scores = scores[np.arange(num_train), y]
92 scores_row_exp_sum = np.sum(np.exp(scores), axis = 1)
93 loss = np.sum(-np.log(np.exp(correct_scores) / scores_row_exp_sum))
94
95 loss /= num_train
96 loss += reg * np.sum(W * W)
97
98
99 # Calculate the gradient of W
100 margins_for_dW = np.exp(scores) / scores_row_exp_sum.reshape(num_train, 1)
101 # The diagonal elements should minus 1
102 margins_for_dW[np.arange(num_train), y] -= 1
103 dW = X.T.dot(margins_for_dW)
104
105 dW /= num_train
106 dW += 2 * reg * W
107
108
109
110
111
112
113
114 # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
115
116 return loss, dW

```

5 neural_net.py

```
1 from __future__ import print_function
2
3 from builtins import range
4 from builtins import object
5 import numpy as np
6 import matplotlib.pyplot as plt
7 from past.builtins import xrange
8
9 class TwoLayerNet(object):
10     """
11     A two-layer fully-connected neural network. The net has an input dimension of
12     N, a hidden layer dimension of H, and performs classification over C classes.
13     We train the network with a softmax loss function and L2 regularization on the
14     weight matrices. The network uses a ReLU nonlinearity after the first fully
15     connected layer.
16
17     In other words, the network has the following architecture:
18
19     input - fully connected layer - ReLU - fully connected layer - softmax
20
21     The outputs of the second fully-connected layer are the scores for each class.
22     """
23
24     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
25         """
26         Initialize the model. Weights are initialized to small random values and
27         biases are initialized to zero. Weights and biases are stored in the
28         variable self.params, which is a dictionary with the following keys:
29
30         W1: First layer weights; has shape (D, H)
31         b1: First layer biases; has shape (H,)
32         W2: Second layer weights; has shape (H, C)
33         b2: Second layer biases; has shape (C,)
34
35         Inputs:
36         - input_size: The dimension D of the input data.
37         - hidden_size: The number of neurons H in the hidden layer.
38         - output_size: The number of classes C.
39         """
40         self.params = {}
41         self.params['W1'] = std * np.random.randn(input_size, hidden_size)
42         self.params['b1'] = np.zeros(hidden_size)
43         self.params['W2'] = std * np.random.randn(hidden_size, output_size)
44         self.params['b2'] = np.zeros(output_size)
45
46     def loss(self, X, y=None, reg=0.0):
47         """
48         Compute the loss and gradients for a two layer fully connected neural
49         network.
50
51         Inputs:
52         - X: Input data of shape (N, D). Each X[i] is a training sample.
53         - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
54             an integer in the range 0 <= y[i] < C. This parameter is optional; if it
55             is not passed then we only return scores, and if it is passed then we
56             instead return the loss and gradients.
57         - reg: Regularization strength.
58
59         Returns:
60         If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
61         the score for class c on input X[i].
62
63         If y is not None, instead return a tuple of:
64         - loss: Loss (data loss and regularization loss) for this batch of training
65             samples.
66         - grads: Dictionary mapping parameter names to gradients of those parameters
67             with respect to the loss function; has the same keys as self.params.
68         """
69         # Unpack variables from the params dictionary
70         W1, b1 = self.params['W1'], self.params['b1']
71         W2, b2 = self.params['W2'], self.params['b2']
72         N, D = X.shape
73
```

```

74 # Compute the forward pass
75 scores = None
76 #####
77 # TODO: Perform the forward pass, computing the class scores for the input. #
78 # Store the result in the scores variable, which should be an array of #
79 # shape (N, C). #
80 #####
81 # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
82
83 h1 = np.maximum(0, np.dot(X, W1) + b1)
84 h2 = np.dot(h1, W2) + b2
85 scores = h2
86
87 # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
88
89 # If the targets are not given then jump out, we're done
90 if y is None:
91     return scores
92
93 # Compute the loss
94 loss = None
95 #####
96 # TODO: Finish the forward pass, and compute the loss. This should include #
97 # both the data loss and L2 regularization for W1 and W2. Store the result #
98 # in the variable loss, which should be a scalar. Use the Softmax #
99 # classifier loss. #
100 #####
101 # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
102
103 num_train = X.shape[0]
104
105 # Avoid numerical instability
106 scores -= np.amax(scores, axis = 1).reshape(scores.shape[0], 1)
107
108 scores_exp = np.exp(scores)
109 scores_exp_sums = np.sum(scores_exp, axis=1)
110 correct_scores_exp = scores_exp[range(num_train), y]
111 loss = correct_scores_exp / scores_exp_sums
112
113 loss = np.sum(-np.log(loss))
114 loss /= num_train
115 loss += reg * (np.sum(W1 * W1) + np.sum(W2 * W2))
116
117 # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
118
119 # Backward pass: compute gradients
120 grads = {}
121 #####
122 # TODO: Compute the backward pass, computing the derivatives of the weights #
123 # and biases. Store the results in the grads dictionary. For example, #
124 # grads['W1'] should store the gradient on W1, and be a matrix of same size #
125 #####
126 # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
127
128 margins_for_dW2 = scores_exp / scores_exp_sums.reshape(num_train, 1)
129
130 # The diagonal elements are correct and they should minus 1
131 margins_for_dW2[np.arange(num_train), y] -= 1
132 margins_for_dW2 /= num_train
133
134 dW2 = h1.T.dot(margins_for_dW2)
135 db2 = np.sum(margins_for_dW2, axis = 0)
136
137 margins_for_hidden = margins_for_dW2.dot(W2.T)
138 margins_for_hidden[h1 <= 0] = 0
139
140 dW1 = X.T.dot(margins_for_hidden)
141 db1 = np.sum(margins_for_hidden, axis = 0)
142
143 dW2 += 2 * reg * W2
144 dW1 += 2 * reg * W1
145
146 grads['W1'] = dW1
147 grads['b1'] = db1
148 grads['W2'] = dW2
149 grads['b2'] = db2

```

```

150 # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
151
152
153 return loss, grads
154
155 def train(self, X, y, X_val, y_val,
156         learning_rate=1e-3, learning_rate_decay=0.95,
157         reg=5e-6, num_iters=100,
158         batch_size=200, verbose=False):
159     """
160     Train this neural network using stochastic gradient descent.
161
162     Inputs:
163     - X: A numpy array of shape (N, D) giving training data.
164     - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
165         X[i] has label c, where 0 ≤ c < C.
166     - X_val: A numpy array of shape (N_val, D) giving validation data.
167     - y_val: A numpy array of shape (N_val,) giving validation labels.
168     - learning_rate: Scalar giving learning rate for optimization.
169     - learning_rate_decay: Scalar giving factor used to decay the learning rate
170         after each epoch.
171     - reg: Scalar giving regularization strength.
172     - num_iters: Number of steps to take when optimizing.
173     - batch_size: Number of training examples to use per step.
174     - verbose: boolean; if true print progress during optimization.
175     """
176     num_train = X.shape[0]
177     iterations_per_epoch = max(num_train / batch_size, 1)
178
179     # Use SGD to optimize the parameters in self.model
180     loss_history = []
181     train_acc_history = []
182     val_acc_history = []
183
184     for it in range(num_iters):
185         X_batch = None
186         y_batch = None
187
188         #####
189         # TODO: Create a random minibatch of training data and labels, storing #
190         # them in X_batch and y_batch respectively. #
191         #####
192         # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
193
194         batch_indices = np.random.choice(num_train, batch_size)
195         X_batch = X[batch_indices]
196         y_batch = y[batch_indices]
197
198         # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
199
200         # Compute loss and gradients using the current minibatch
201         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
202         loss_history.append(loss)
203
204         #####
205         # TODO: Use the gradients in the grads dictionary to update the #
206         # parameters of the network (stored in the dictionary self.params) #
207         # using stochastic gradient descent. You'll need to use the gradients #
208         # stored in the grads dictionary defined above. #
209         #####
210         # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
211
212         self.params['W1'] += -grads['W1']*learning_rate
213         self.params['b1'] += -grads['b1']*learning_rate
214         self.params['W2'] += -grads['W2']*learning_rate
215         self.params['b2'] += -grads['b2']*learning_rate
216
217         # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
218
219         if verbose and it % 100 == 0:
220             print('iteration %d / %d: loss %f' % (it, num_iters, loss))
221
222         # Every epoch, check train and val accuracy and decay learning rate.
223         if it % iterations_per_epoch == 0:
224             # Check accuracy
225             train_acc = (self.predict(X_batch) == y_batch).mean()

```

```

226         val_acc = (self.predict(X_val) == y_val).mean()
227         train_acc_history.append(train_acc)
228         val_acc_history.append(val_acc)
229
230         # Decay learning rate
231         learning_rate *= learning_rate_decay
232
233     return {
234         'loss_history': loss_history,
235         'train_acc_history': train_acc_history,
236         'val_acc_history': val_acc_history,
237     }
238
239 def predict(self, X):
240     """
241     Use the trained weights of this two-layer network to predict labels for
242     data points. For each data point we predict scores for each of the C
243     classes, and assign each data point to the class with the highest score.
244
245     Inputs:
246     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
247         classify.
248
249     Returns:
250     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
251         the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
252         to have class c, where 0 <= c < C.
253     """
254     y_pred = None
255
256     #####
257     # TODO: Implement this function; it should be VERY simple! #
258     #####
259     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
260
261     W1, b1 = self.params['W1'], self.params['b1']
262     W2, b2 = self.params['W2'], self.params['b2']
263     h1 = np.maximum(0, np.dot(X, W1) + b1)
264     h2 = np.dot(h1, W2) + b2
265     scores = h2
266     y_pred = np.argmax(scores, axis = 1)
267
268     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
269
270     return y_pred

```